

Module 2 - Session 1:

Data and Model Building

Module 2 Overview

What will we learn?

- Data Pipeline → handle large datasets
- Beyond Sequential → custom architectures
- Optimizers
- Device Management
- Building an image classifier



Session 1: Data and Model Building

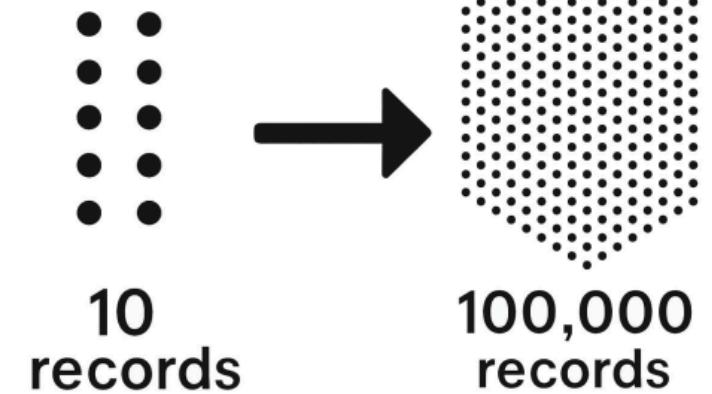
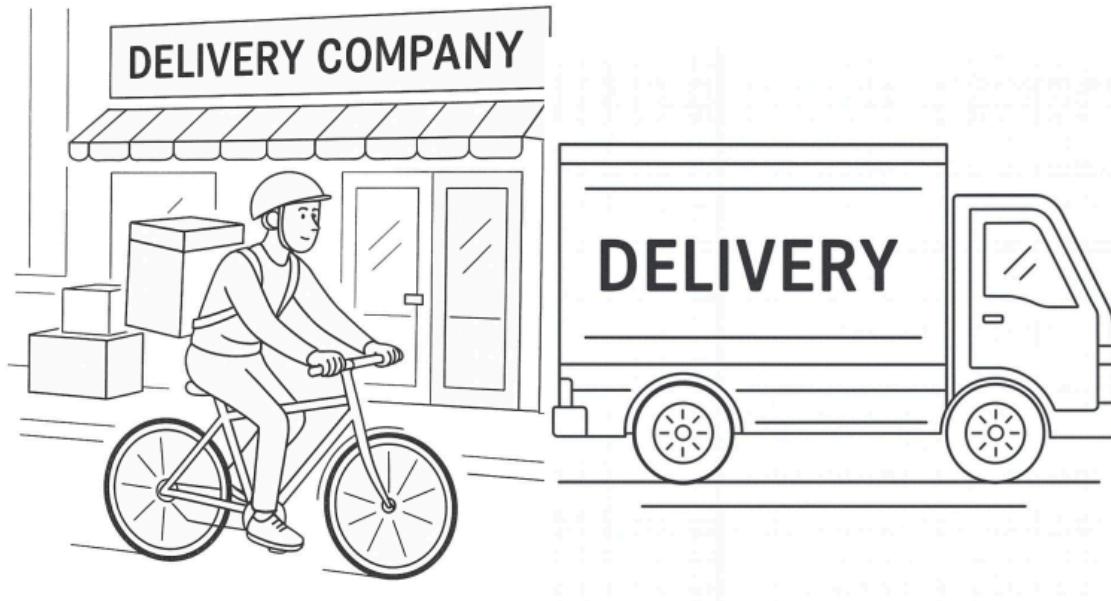
- Revisit the ML pipeline with a focus on PyTorch's data handling tools
- Learn about data management at scale
- Explore building custom model architectures beyond the Sequential API



The Challenge: Large Datasets

100,000 delivery records

Problem: Loading all at once → runs out of memory



Solution: Work with data in batches

PyTorch Data Utilities

Three core tools:

1. **Transforms** - operations on each data point
2. **Dataset** - fetches samples from disk on demand
3. **DataLoader** - serves data in batches



1. Transforms

```
1 transform = transforms.Compose([  
2     transforms.ToTensor(),  
3     transforms.Normalize(mean=0.5, std=0.5)  
4 ])
```



ToTensor: converts to tensors, scales 0-255 → 0-1

Normalize: centers around 0, scales using standard deviation

2. Dataset

```
1 dataset = MNIST(root='./data', train=True,  
2                   download=True, transform=transform)
```

Key features:

- Fetches samples from disk when asked
- Doesn't preload everything
- Handles where data lives, how to load samples, total count

3. DataLoader

```
1 dataloader = DataLoader(dataset, batch_size=64, shuffle=True)
```



Batch size: how many samples per batch

Shuffle: randomize order each epoch

Makes training on large datasets possible



Complete Data Pipeline

```
1 # 1. Define transforms
2 transform = transforms.Compose([...])
3
4 # 2. Create dataset
5 train_dataset = MNIST(..., transform=transform)
6
7 # 3. Create dataloader
8 train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
9
10 # 4. Use in training loop
11 for batch in train_loader:
12     images, labels = batch
13     # train model
```



Beyond Sequential: Custom Models

nn.Sequential

```
1 model = nn.Sequential(    □  
2     nn.Linear(1, 20),  
3     nn.ReLU(),  
4     nn.Linear(20, 1)  
5 )
```

nn.Module

```
1 class MyModel(nn.Module):    □  
2     # defines layers  
3     def __init__(self):  
4         super().__init__()  
5         self.layer1 = nn.Linear(1, 20)  
6         self.layer2 = nn.Linear(20, 1)  
7  
8     # describes data flow  
9     def forward(self, x):  
10        x = self.layer1(x)  
11        x = F.relu(x)  
12        x = self.layer2(x)  
13        return x
```

More control, same functionality

Calling the Model

Don't call `model.forward()` directly

Do call `model(input)`

PyTorch handles the forward call and essential bookkeeping

Why `super().__init__()`?

Necessary for parameter tracking

PyTorch needs to set up a system to track all learnable parameters (weights and biases)

Without it, PyTorch has nowhere to register your layers

Training Loop Pattern

```
1 for batch in dataloader:  
2     optimizer.zero_grad()      # Clear old calculations  
3     outputs = model(inputs)    # Forward pass  
4     loss = loss_fn(outputs, targets) # Measure error  
5     loss.backward()            # Calculate gradients  
6     optimizer.step()          # Update weights
```



Order matters! Don't swap these steps.

Evaluation

```
1 model.eval() # Set to evaluation mode
2 with torch.no_grad(): # Disable gradient tracking
3     for batch in test_loader:
4         outputs = model(inputs)
5         # Calculate accuracy
```



Two critical things:

- `model.eval()` - sets evaluation mode
- `torch.no_grad()` - disables gradient tracking

Measuring Performance

For classification: Accuracy

```
1 correct = (predictions == labels).sum().item()  
2 total = labels.size(0)  
3 accuracy = correct / total
```

Count correct predictions / total predictions

To sum up

- Data pipeline: Dataset, DataLoader, Transforms
- Model building: `nn.Module`
- Training loop: `for batch in dataloader:`
- Evaluation: `model.eval()`, `torch.no_grad()`, `accuracy`

What's Next?

In **Session 2: Loss Functions and Optimizers**, you learn:

- How loss functions measure error
- Cross-entropy loss for classification
- How optimizers use gradients to update weights
- Understanding backpropagation