

Module 2 - Session 3:

Device Management

and Image Classification



Setup



Session 3



Device Management



Every tensor and model lives on a device

- CPU (default)
- GPU (accelerator)

Key rule: Model and data must be on the same device!

CPU vs GPU

CPU:

- Default device
- General purpose
- Sequential operations

GPU:

- Accelerator
- Parallel operations
- **10-15x faster for training**

Checking for GPU

```
1 torch.cuda.is_available() # Returns True if GPU available
```



Common pattern:

```
1 device = torch.device('cuda' if torch.cuda.is_available()
2                           else 'cpu')
```



Moving Model to Device

```
1 model = MyModel()  
2 model = model.to(device) # Move model to device
```



Puts model's parameters on the selected device

Moving Data to Device

```
1 for batch in dataloader:  
2     inputs, targets = batch  
3     inputs = inputs.to(device)  
4     targets = targets.to(device)  
5     # ... rest of training
```



Move each batch within the training loop

Checking Device Location

```
1 # For tensors
2 tensor.device
3
4 # For models (check a parameter)
5 next(model.parameters()).device
```



Common Mistake with `.to()`

.to() doesn't change tensor in place

It creates a new one!

```
1 # Wrong
2 tensor.to(device) # Result is discarded!
3
4 # Right
5 tensor = tensor.to(device) # Reassign
```



Complete Training Loop with Device Management

```
1 device = torch.device('cuda' if torch.cuda.is_available()
2                         else 'cpu')
3 model = MyModel().to(device)
4
5 for batch in dataloader:
6     inputs, targets = batch[0].to(device), batch[1].to(device)
7     optimizer.zero_grad()
8     outputs = model(inputs)
9     loss = loss_fn(outputs, targets)
10    loss.backward()
11    optimizer.step()
```

Key steps: 1. Choose device up front 2. Move model once 3. Move data in every batch

GPU Memory Limits

GPU memory is limited

Error if batch size too large:


RuntimeError: CUDA out of memory

Solution: Lower batch size (32-64 is good starting point)



Building Your First Image Classifier

28x28 Image



(a) MNIST sample belonging to the digit '7'.

Dataset: 60,000 images

0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9

(b) 100 samples from the MNIST training set.

MNIST Dataset:

- 60,000 training images
- 10,000 test images
- 28×28 pixels, grayscale
- 10 classes (digits 0-9)

Setting Up the Data Pipeline

```
1 import torchvision  
2 from torchvision import transforms  
3  
4 transform = transforms.Compose([  
5     transforms.ToTensor(),  
6     transforms.Normalize((0.1307,), (0.3081,))  
7 ])
```

ToTensor: converts to tensors, scales 0-255 → 0-1

Normalize: centers around 0 using dataset mean/std

Loading the Dataset

```
1 train_dataset = torchvision.datasets.MNIST(  
2     root='./data',  
3     train=True,  
4     download=True,  
5     transform=transform  
6 )  
7  
8 test_dataset = torchvision.datasets.MNIST(  
9     root='./data',  
10    train=False,  
11    transform=transform  
12 )
```



TorchVision handles downloading and organizing

Creating DataLoaders

```
1 train_loader = DataLoader(  
2     train_dataset,  
3     batch_size=64,  
4     shuffle=True  
5 )  
6  
7 test_loader = DataLoader(  
8     test_dataset,  
9     batch_size=1000,  
10    shuffle=False  
11 )
```

Training: shuffle=True (mix up order each epoch)

Testing: shuffle=False (order doesn't matter)

Building the Model Architecture

```
1 class MNISTClassifier(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.flatten = nn.Flatten()
5         self.linear1 = nn.Linear(784, 128)
6         self.relu = nn.ReLU()
7         self.linear2 = nn.Linear(128, 10)
8
9     def forward(self, x):
10        x = self.flatten(x)
11        x = self.linear1(x)
12        x = self.relu(x)
13        x = self.linear2(x)
14        return x
```




Why Flatten?

MNIST images: shape [1, 28, 28] (channels, height, width)

With batch: shape [64, 1, 28, 28]

Linear layers expect: flat vectors [batch, features]

Flatten: [64, 1, 28, 28] \rightarrow [64, 784]



Model Architecture Breakdown

Linear(784, 128):

- 784 pixel values → 128 hidden features

ReLU:

- Activation function (non-linearity)

Linear(128, 10):

- 128 features → 10 outputs (one per digit class)

What's Next?

In **Session 4: Training and Evaluating Your Classifier** we learn:

- Setting up loss function and optimizer
- Writing the training loop
- Evaluating on test set
- Watching your model learn!