

Module 2 - Session 4: Training and Evaluating Your Classifier

Session 4



Setting Up Training

Everything you need:

- Model (on device)
- Loss function
- Optimizer
- Data loaders

Device and Model Setup

```
1 device = torch.device('cuda' if torch.cuda.is_available()
2                         else 'cpu')
3 model = MNISTClassifier().to(device)
```

Both model and data must be on the same device



Loss Function

```
1 loss_function = nn.CrossEntropyLoss()
```



Designed for classification tasks

Perfect for choosing a digit from 0-9



Optimizer

```
1 optimizer = optim.Adam(model.parameters(), lr=0.001)
```



Adam: adapts learning rate as it trains

- Larger adjustments early (noisy gradients)
- Smaller corrections later (training stabilizes)



Training Function

```
1 def train_one_epoch(model, dataloader, loss_fn, optimizer, device):
2     model.train() # Set to training mode
3     running_loss = 0.0
4     correct = 0
5     total = 0
6
7     for batch_idx, (data, targets) in enumerate(dataloader):
8         # Move to device
9         data, targets = data.to(device), targets.to(device)
10
11     # Training steps
12     optimizer.zero_grad()
13     outputs = model(data)
14     loss = loss_fn(outputs, targets)
15     loss.backward()
16     optimizer.step()
17
18     # Track progress
19     running_loss += loss.item()
20     _, predicted = outputs.max(1)
21     total += targets.size(0)
22     correct += predicted.eq(targets).sum().item()
23
24     if batch_idx % 100 == 0:
25         print(f'Loss: {loss.item():.4f}, '
26               f'{batch_idx * len(data)} / {len(dataloader.dataset)}')
```



Understanding the Training Progress

With 60,000 training images and batch size 64:

- About 938 batches per epoch
- Around 9 progress updates

Watch the numbers:

- Loss: dropping ($0.64 \rightarrow 0.17$)
- Accuracy: climbing ($81\% \rightarrow 95\%$)



Evaluation Function

```
1 def evaluate(model, dataloader, device):
2     model.eval()    # Set to evaluation mode
3     correct = 0
4     total = 0
5
6     with torch.no_grad():    # Disable gradient tracking
7         for data, targets in dataloader:
8             data, targets = data.to(device), targets.to(device)
9             outputs = model(data)
10            _, predicted = outputs.max(1)
11            total += targets.size(0)
12            correct += predicted.eq(targets).sum().item()
13
14    return 100. * correct / total
```



Putting It All Together

```
1 num_epochs = 10
2
3 for epoch in range(num_epochs):
4     print(f'Epoch {epoch+1}/{num_epochs}')
5
6     # Train
7     train_one_epoch(model, train_loader, loss_function,
8                      optimizer, device)
9
10    # Evaluate
11    accuracy = evaluate(model, test_loader, device)
12    print(f'Test Accuracy: {accuracy:.2f}%')
13    print('-' * 50)
```

10 epochs = 10 full passes through training data

After each epoch: evaluate on test set

What You'll See

Epoch: 1

| | | |
|-----------------------|--|-------------|
| Loss: 0.627 | | Acc: 81.97% |
| Loss: 0.336 | | Acc: 90.28% |
| Loss: 0.251 | | Acc: 92.44% |
| Loss: 0.235 | | Acc: 93.31% |
| Loss: 0.213 | | Acc: 93.64% |
| Loss: 0.190 | | Acc: 94.66% |
| Loss: 0.167 | | Acc: 95.12% |
| Loss: 0.159 | | Acc: 95.34% |
| Loss: 0.155 | | Acc: 95.52% |
| Test Accuracy: 95.48% | | |

Epoch: 2

| | | |
|-----------------------|--|-------------|
| Loss: 0.115 | | Acc: 96.78% |
| Loss: 0.112 | | Acc: 96.66% |
| Loss: 0.122 | | Acc: 96.39% |
| Loss: 0.116 | | Acc: 96.59% |
| Loss: 0.105 | | Acc: 96.91% |
| Loss: 0.104 | | Acc: 96.78% |
| Loss: 0.116 | | Acc: 96.48% |
| Loss: 0.105 | | Acc: 96.64% |
| Loss: 0.095 | | Acc: 97.05% |
| Test Accuracy: 97.02% | | |

Epoch: 10

| | | |
|-----------------------|--|-------------|
| Loss: 0.013 | | Acc: 99.69% |
| Loss: 0.015 | | Acc: 99.50% |
| Loss: 0.017 | | Acc: 99.45% |
| Loss: 0.019 | | Acc: 99.41% |
| Loss: 0.026 | | Acc: 99.08% |
| Loss: 0.022 | | Acc: 99.11% |
| Loss: 0.028 | | Acc: 99.00% |
| Loss: 0.028 | | Acc: 99.02% |
| Loss: 0.023 | | Acc: 99.22% |
| Test Accuracy: 97.58% | | |

By epoch 10: Loss: tiny, Accuracy: high (often 95%+)

When accuracy stops improving: Model is done learning, May not need all 10 epochs

Module 2 Summary

You've learned:

- PyTorch data pipeline (Dataset, DataLoader, Transforms)
- Building custom models with `nn.Module`
- Loss functions (MSE vs Cross-Entropy)
- How optimizers use gradients to update weights
- Device management (CPU vs GPU)
- Complete image classification pipeline

Lab 1: Building Your First Image Classifier

“Don’t tell me the moon is shining; show me the glint of light on broken glass.”

CUE: START THE LAB HERE

Assignment 2: EMNIST Letter Detective

“The eye sees only what the mind is prepared to comprehend.”

CUE: START THE ASSIGNMENT HERE

What's Next?

In **Module 3: Data Management** we learn:

- How to manage and preprocess data for deep learning
- How to build a robust data pipeline
- How to use data augmentation to improve model performance
- How to use data validation to ensure model performance